

```
In [1]: import pandas as pd
```

```
In [2]: import numpy as np
```

```
In [12]: import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [3]: df=pd.read_csv(r'C:\Users\Rutu\Documents\New folder\diabetes.csv')
```

```
In [4]: df.head()
```

```
Out[4]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Pedigree	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

```
In [5]: df.tail()
```

```
Out[5]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Pedigree	Age	Outcome
763	10	101	76	48	180	32.9	0.171	63	0
764	2	122	70	27	0	36.8	0.340	27	0
765	5	121	72	23	112	26.2	0.245	30	0
766	1	126	60	0	0	30.1	0.349	47	1
767	1	93	70	31	0	30.4	0.315	23	0

```
In [6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
 #   Column              Non-Null Count  Dtype  
---  -
 0   Pregnancies         768 non-null   int64  
 1   Glucose              768 non-null   int64  
 2   BloodPressure        768 non-null   int64  
 3   SkinThickness        768 non-null   int64  
 4   Insulin              768 non-null   int64  
 5   BMI                  768 non-null   float64 
 6   Pedigree             768 non-null   float64 
 7   Age                  768 non-null   int64  
 8   Outcome              768 non-null   int64  
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
In [7]: df.isnull().sum()
```

```
Out[7]: Pregnancies    0
Glucose              0
BloodPressure        0
SkinThickness        0
```

```
Insulin      0
BMI           0
Pedigree      0
Age           0
Outcome      0
dtype: int64
```

```
In [9]: df.shape
```

```
Out[9]: (768, 9)
```

```
In [10]: df.describe()
```


	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Pedigree	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	0.471876	33.240885	0.481561
std	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	0.331329	11.760232	0.500141
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.078000	21.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	0.243750	24.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	0.372500	29.000000	0.000000
75%	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	0.626250	41.000000	0.000000
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	2.420000	81.000000	0.000000

```
In [11]: df.corr()
```


	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Pedigree	Age	Outcome
Pregnancies	1.000000	0.129459	0.141282	-0.081672	-0.073535	0.017683	-0.033523	0.544341	0.281837
Glucose	0.129459	1.000000	0.152590	0.057328	0.331357	0.221071	0.137337	0.263514	0.481561
BloodPressure	0.141282	0.152590	1.000000	0.207371	0.088933	0.281805	0.041265	0.239528	0.000000
SkinThickness	-0.081672	0.057328	0.207371	1.000000	0.436783	0.392573	0.183928	-0.113970	0.000000
Insulin	-0.073535	0.331357	0.088933	0.436783	1.000000	0.197859	0.185071	-0.042163	0.000000
BMI	0.017683	0.221071	0.281805	0.392573	0.197859	1.000000	0.140647	0.036242	0.281837
Pedigree	-0.033523	0.137337	0.041265	0.183928	0.185071	0.140647	1.000000	0.033561	0.000000
Age	0.544341	0.263514	0.239528	-0.113970	-0.042163	0.036242	0.033561	1.000000	0.281837
Outcome	0.281837	0.481561	0.000000	0.000000	0.000000	0.281837	0.000000	0.281837	1.000000

```
In [14]: df['Glucose'].fillna(df['Glucose'].mean(), inplace = True)
df['BloodPressure'].fillna(df['BloodPressure'].mean(), inplace = True)
df['SkinThickness'].fillna(df['SkinThickness'].median(), inplace = True)
df['Insulin'].fillna(df['Insulin'].median(), inplace = True)
df['BMI'].fillna(df['BMI'].median(), inplace = True)
```

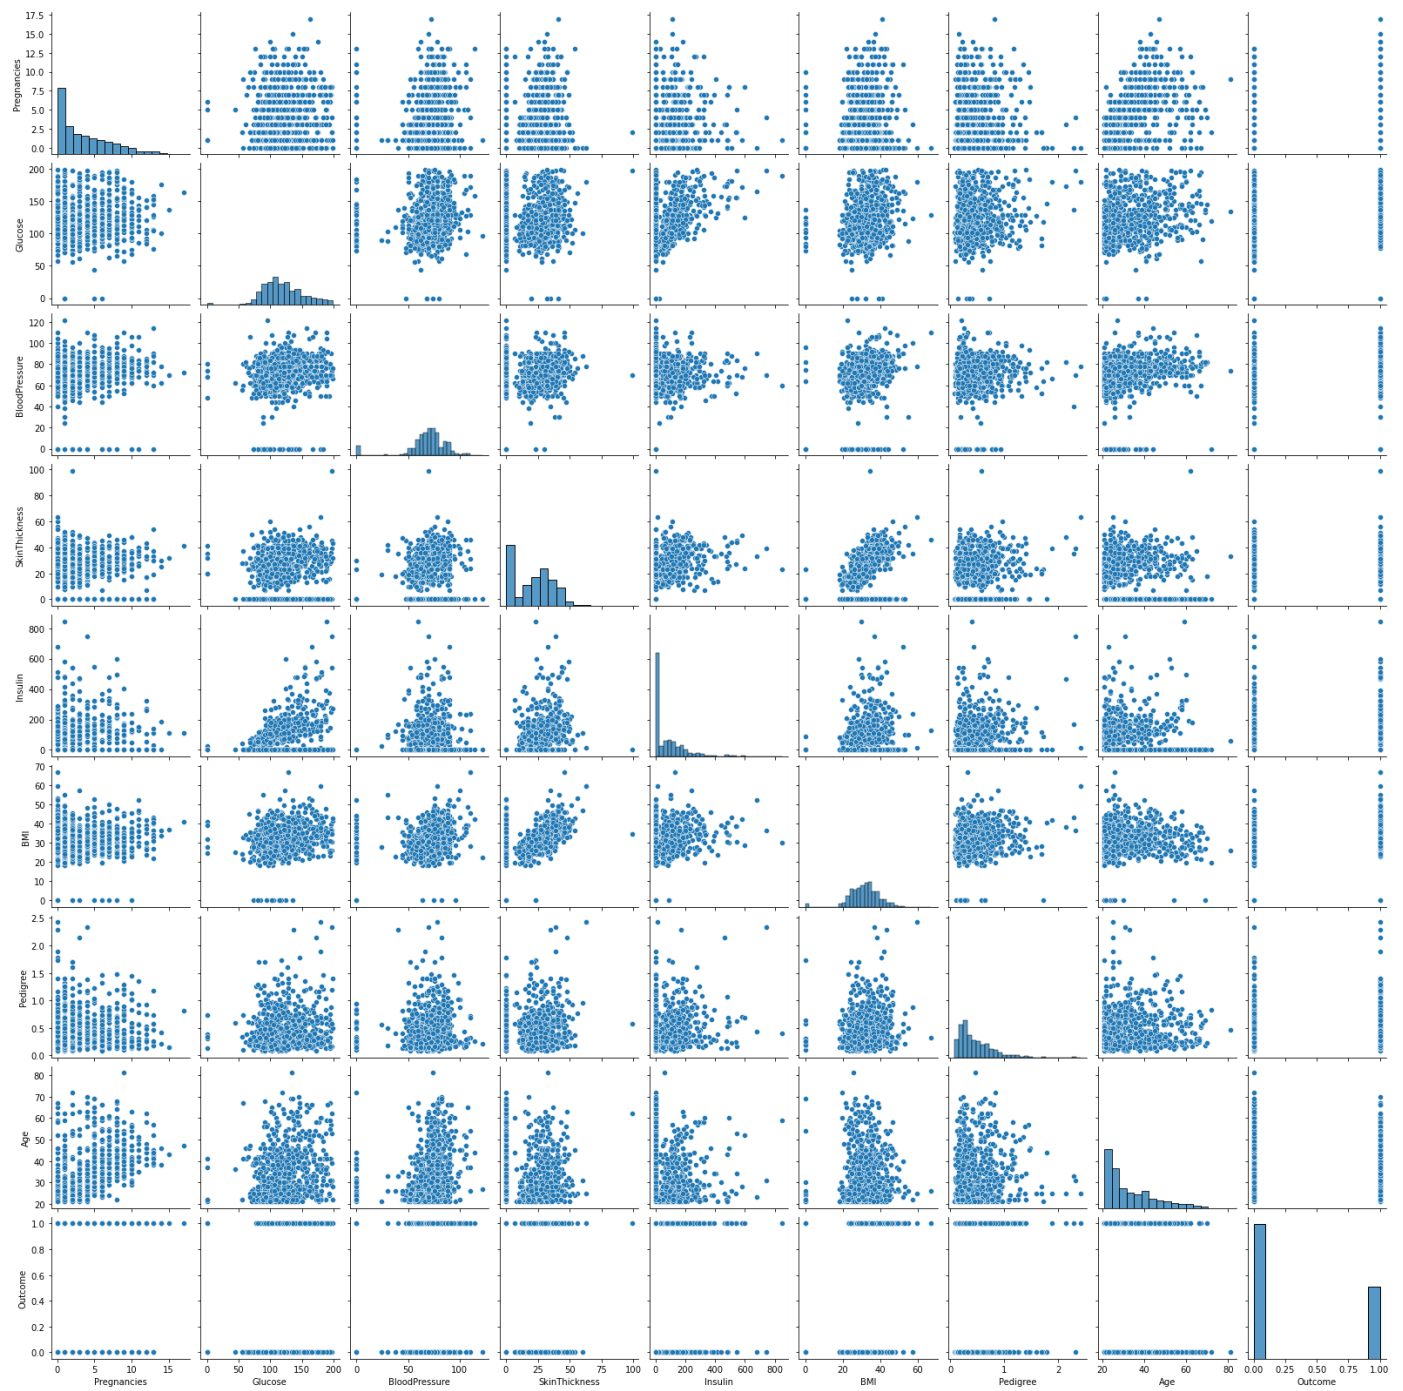
```
In [16]: p = df.hist(figsize = (20,20))
```




```
In [19]: sns.countplot(y=df.dtypes ,data=df)
plt.xlabel("count of each data type")
plt.ylabel("data types")
plt.show()
```



```
In [23]: p=sns.pairplot(df)
```



```
In [25]: plt.figure(figsize=(12,10)) # on this line I just set the size of figure to 12 by 10.
p=sns.heatmap(df.corr(), annot=True,cmap='RdYlGn')
```



```
In [26]: from sklearn.preprocessing import StandardScaler
sc_X = StandardScaler()
X = pd.DataFrame(sc_X.fit_transform(df.drop(["Outcome"],axis = 1)),
                  columns=['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
                           'BMI', 'DiabetesPedigreeFunction', 'Age'])
```

```
In [30]: y = df.Outcome
```

```
In [28]: X.head()
```

```
Out[28]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age
0	0.639947	0.848324	0.149641	0.907270	-0.692891	0.204013	0.468492	1.425995
1	-0.844885	-1.123396	-0.160546	0.530902	-0.692891	-0.684422	-0.365061	-0.190672
2	1.233880	1.943724	-0.263941	-1.288212	-0.692891	-1.103255	0.604397	-0.105584
3	-0.844885	-0.998208	-0.160546	0.154533	0.123302	-0.494043	-0.920763	-1.041545
4	-1.141852	0.504055	-1.504687	0.907270	0.765836	1.409746	5.484909	-0.020496

```
In [31]: from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=1/3,random_state=42, stra
```

In [32]: `from sklearn.neighbors import KNeighborsClassifier`

```
test_scores = []
train_scores = []

for i in range(1,15):

    knn = KNeighborsClassifier(i)
    knn.fit(X_train,y_train)

    train_scores.append(knn.score(X_train,y_train))
    test_scores.append(knn.score(X_test,y_test))
```

In [33]: `max_train_score = max(train_scores)`
`train_scores_ind = [i for i, v in enumerate(train_scores) if v == max_train_score]`
`print('Max train score {} % and k = {}'.format(max_train_score*100,list(map(lambda x: x+`
`Max train score 100.0 % and k = [1]`

In [34]: `max_test_score = max(test_scores)`
`test_scores_ind = [i for i, v in enumerate(test_scores) if v == max_test_score]`
`print('Max test score {} % and k = {}'.format(max_test_score*100,list(map(lambda x: x+1,`
`Max test score 76.953125 % and k = [11]`

In [35]: `knn = KNeighborsClassifier(11)`
`knn.fit(X_train,y_train)`
`knn.score(X_test,y_test)`

Out[35]: 0.76953125

In [39]: `from sklearn.metrics import confusion_matrix`
#let us get the predictions using the classifier we had fit above
`y_pred = knn.predict(X_test)`
`confusion_matrix(y_test,y_pred)`
`pd.crosstab(y_test, y_pred, rownames=['True'], colnames=['Predicted'], margins=True)`


Out[39]:

	Predicted	0	1	All
True				
0	147	20	167	
1	39	50	89	
All	186	70	256	

In [40]: `y_pred = knn.predict(X_test)`
`from sklearn import metrics`
`cnf_matrix = metrics.confusion_matrix(y_test, y_pred)`
`p = sns.heatmap(pd.DataFrame(cnf_matrix), annot=True, cmap="YlGnBu" ,fmt='g')`
`plt.title('Confusion matrix', y=1.1)`
`plt.ylabel('Actual label')`
`plt.xlabel('Predicted label')`

Out[40]: Text(0.5, 15.0, 'Predicted label')

Confusion matrix



```
In [46]: from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.79	0.88	0.83	167
1	0.71	0.56	0.63	89
accuracy			0.77	256
macro avg	0.75	0.72	0.73	256
weighted avg	0.76	0.77	0.76	256

```
In [ ]:
```